

DeepMind

# Graph Neural Networks in Computational Biology

*(a Personal Perspective)*

Petar Veličković

Computational Biology Society Seminar  
Imperial College London  
19 April 2021



DeepMind


In this talk:  
**Graph neural networks for biological data**




# What this talk is ***not!***



# What this talk is *not!* 🙄



T1037 / 6vr4  
90.7 GDT  
(RNA polymerase domain)



T1049 / 6y4f  
93.3 GDT  
(adhesin tip)

- Experimental result
- Computational prediction

**AlphaFold: a solution to a 50-year-old grand challenge in biology**






For more on AlphaFold, see:

[https://deepmind.com/blog/article/alphafold-a-solution-to-a-50-year-old-gr  
and-challenge-in-biology](https://deepmind.com/blog/article/alphafold-a-solution-to-a-50-year-old-grand-challenge-in-biology)



# What this talk **is**

- Hopefully an exciting field from *many angles* :)
  - Molecular interactions
  - Protein function prediction
  - Genome assembly
  - Computational neuroscience
  - Electronic Health Records



# What this talk **is**

- Hopefully an exciting field from *many* angles :)
  - Molecular interactions
  - Protein function prediction
  - Genome assembly
  - Computational neuroscience
  - Electronic Health Records
- More broadly...
  - Personal perspective on this rich, interdisciplinary field
  - For ML audience: **you can do it!**
    - + a **blueprint** for approaching the area
  - For Bio audience: hopefully a useful **computational tool**  
(for both: interdisciplinary collaboration can work wonders!)



# Let's start at the beginning

- Born in Belgrade (Serbia) in the 1990s
  - Family members worked for local representatives of “big pharma” (Merck)



# Let's start at the beginning

- Born in Belgrade (Serbia) in the 1990s
  - Family members worked for local representatives of “big pharma” (Merck)
- Gradually increasing interest towards **computer science** – especially classical algorithms



# Let's start at the beginning

- Born in Belgrade (Serbia) in the 1990s
  - Family members worked for local representatives of “big pharma” (Merck)
- Gradually increasing interest towards **computer science** – especially classical algorithms
- Developed strong interest in biology in high school (primarily thanks to **Branka Dobrković**)

I'm a biologist, and Petar already has high experience in computer science. To me this combination seems like an ideal link for very attractive scientific disciplines in the world – bioinformatics and similar. It seems to me like this connection of natural sciences with computer science would be the perfect choice for him.



# Let's start at the beginning

- Born in Belgrade (Serbia) in the 1990s
  - Family members worked for local representatives of “big pharma” (Merck)
- Gradually increasing interest towards **computer science** – especially classical algorithms
- Developed strong interest in biology in high school (primarily thanks to **Branka Dobrković**)
- Computer Science at Cambridge (2012–15)
  - Lost nearly all contact with biology
- Reached out to Prof Pietro Liò for my final-year project
  - Realised that bioinformatics is **brimming** with classical algorithms
  - Pietro suggested a project in *machine learning*, however...
  - The rest is history (i.e. this talk)



# Before GNNs...

- I started my PhD in 2016, with a paper classifying **breast cancer**
- Officially I was a “Research Assistant in Computational Biology”
  - But **no formal training** in biology!
  - Luckily, the field is remarkably accessible and full of interesting problems to solve
  - It was **very** helpful to talk to domain experts and understand the “burning questions”
- Fruitful collaborations lead to **Parapred** (Bioinformatics) and **ChronoMID** (PLOS ONE)
  - Carefully crafted machine learning solutions to problems posed by domain experts

## Parapred: antibody paratope prediction using convolutional and recurrent neural networks FREE

Edgar Liberis ✉, Petar Veličković, Pietro Sormanni ✉, Michele Vendruscolo, Pietro Liò

Bioinformatics, Volume 34, Issue 17, 01 September 2018, Pages 2944–2950,

<https://doi.org/10.1093/bioinformatics/bty305>

Published: 16 April 2018 Article history ▾

## ChronoMID—Cross-modal neural networks for 3-D temporal medical imaging data

Alexander G. Rakowski, Petar Veličković, Enrico Dall'Ara ✉, Pietro Liò

Published: February 21, 2020 • <https://doi.org/10.1371/journal.pone.0228962>





# Before GNNs...

- I started my PhD in 2016, with a paper classifying **breast cancer**
- Officially I was a “Research Assistant in Computational Biology”
  - But **no formal training** in biology!
  - Luckily, the field is remarkably accessible and full of interesting problems to solve
  - It was **very** helpful to talk to domain experts and understand the “burning questions”
- Fruitful collaborations lead to **Parapred** (Bioinformatics) and **ChronoMID** (PLOS ONE)
  - Carefully crafted machine learning solutions to problems posed by domain experts
- “Game changing” moment in 2017, when I discovered **graph representation learning**
  - Why should you care?




# Molecules are graphs!

- A very natural way to represent molecules is as a **graph**
  - **Atoms** as nodes, **bonds** as edges
  - Features such as **atom type, charge, bond type...**



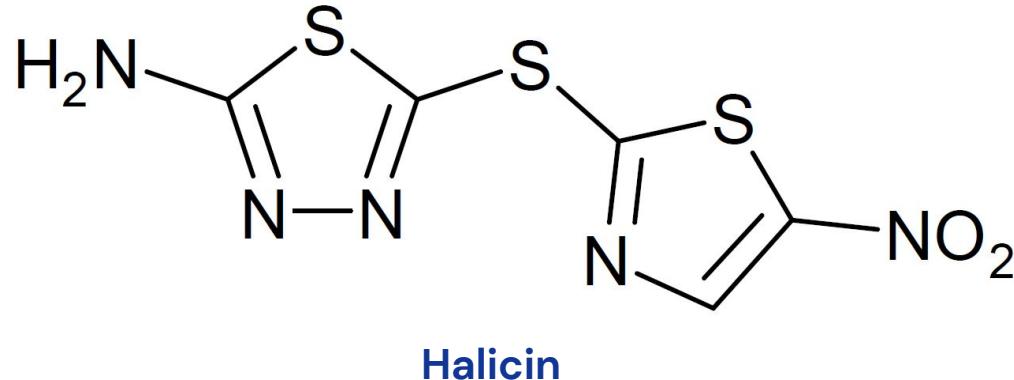
# GNNs for molecule classification

- Interesting task to predict is, for example, whether the molecule is a potent **drug**.
  - Can do *binary classification* on whether the drug will inhibit certain bacteria. (*E.coli*)
  - Train on a **curated dataset** for compounds where response is known.



# Follow-up study

- Once trained, the model can be applied to *any* molecule.
  - Execute on a large dataset of known candidate molecules.
  - Select the ~top-100 candidates from your GNN model.
  - Have chemists thoroughly investigate those (after some additional filtering).
- Discover a previously overlooked compound that is a **highly potent** antibiotic!



# ...Achieve wide acclaim!

Arguably the most popularised **success story** of graph neural networks to date!

**Cell**

**A Deep Learning Approach to Antibiotic Discovery**

**Graphical Abstract**

The graphical abstract illustrates a deep learning pipeline for antibiotic discovery. It starts with a 'Directed message passing neural network' (DMPNN) processing two molecules. The output of the DMPNN is used to predict antibiotic activity ('Antibiotic predictions (upper limit  $10^{11} +$ )'). This prediction is compared against a 'Training set ( $10^4$  molecules)' to validate the model ('Model validation'). The validation process involves 'Growth [antibiotic]' curves. The validated model is then used to predict activity for a large dataset ('Chemical space') represented as a funnel. The final output is 'New Antibiotics'. A legend indicates: 1. Training set; 2. Model validation; 3. Antibiotic predictions; 4. DMPNN; 5. Chemical space.

**Authors**  
Jonathan M. Stokes, Kevin Yang, Kyle Swanson, ..., Tommi S. Jaakkola, Regina Barzilay, James J. Collins

**Correspondence**  
[regina@csail.mit.edu](mailto:regina@csail.mit.edu) (R.B.), [jimjc@mit.edu](mailto:jimjc@mit.edu) (J.J.C.)

**In Brief**  
A trained deep neural network predicts antibiotic activity in molecules that are structurally different from known antibiotics, among which Halicin exhibits efficacy against broad-spectrum bacterial infections in mice.

**Drug Repurposing Hub HALICIN**  
Halicin is shown interacting with a bacterial cell, causing a change in pH ( $\Delta\text{pH}$ ) and leading to bacterial cell death. The chemical structure of Halicin is shown: CN1=CSC(=S)c2cc([N+]([O-])=O)cc21.

**ZINC15 Database**  
The ZINC15 Database contains various chemical compounds. Two specific ones are highlighted:  
1. Rapidly bactericidal Broad-spectrum: O=[N+]([O-])c1cc(Br)c(O)nc2c(N)nc3c(O)nnc4c3[nH]c24  
2. Low MIC Broad-spectrum: N#Cc1ccc(cc1)N2Cc3c(cc(cc3)N(C)c4ccccc4)C(=O)N2Cc5ccccc5

**(Stokes et al., Cell'20)**



...Achieve wide acclaim!

Arguably the most popularised **success story** of graph neural networks to date!

The screenshot shows a news article from the journal **nature**. At the top right is a blue "Subscribe" button. Below the header, the text "ARTICLE" is visible. The main headline reads: **Powerful antibiotics discovered using AI**. Below the headline is a sub-headline: "Machine learning spots molecules that work even against 'untreatable' strains of bacteria." At the bottom left, it says "(Stokes et al., Cell'20)". A small graphic at the bottom features a mouse icon, the names of two bacteria (*Acinetobacter baumannii* and *Clostridioides difficile*), a green checkmark, and a chemical structure labeled "Broad-spectrum".

ARTICLE

**nature**

Subscribe

NEWS · 20 FEBRUARY 2020

# Powerful antibiotics discovered using AI

Machine learning spots molecules that work even against 'untreatable' strains of bacteria.

(Stokes et al., Cell'20)

Acinetobacter baumannii  
Clostridioides difficile

Broad-spectrum



# ...Achieve wide acclaim!

Arguably the most popular

# nature

NEWS · 20 FEBRUARY 2020

## Powerful and

Machine learning spots  
bacteria.

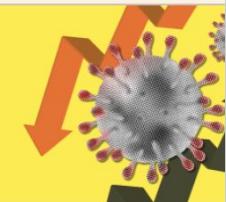
(Stokes et al., Cell'20)

# FINANCIAL TIMES

S COMPANIES TECH MARKETS GRAPHICS OPINION WORK & CAREERS LIFE & ARTS HOW TO SPEND IT

## CORONAVIRUS BUSINESS UPDATE

Get 30 days' complimentary access to our Coronavirus Business Update newsletter



Artificial intelligence

Robotics



'Death of the office' homeworking claims exaggerated



Anti-social robots harr  
increase social distanc

Artificial intelligence

+ Add to myFT

## AI discovers antibiotics to treat drug-resistant diseases

Machine learning uncovers potent new drug able to kill 35 powerful bacteria



...Achieve wide acclaim!

The image displays a screenshot of a BBC News website. At the top, there is a navigation bar with the BBC logo, a 'Sign in' button, and links for 'News', 'Sport', 'Reel', 'Worklife', 'Travel', and 'Future'. Below this is a large red banner with the word 'NEWS' in white. Underneath the banner is a secondary navigation bar with links for 'Home', 'Video', 'World', 'UK', 'Business', 'Tech', 'Science', 'Stories', and 'Entertainment & Arts'. A blue box labeled 'BBC WORKLIFE' is positioned on the left side of the main content area. The main headline reads 'Our new guide for getting ahead'. Below the headline, a large article summary is visible, featuring the title 'Scientists discover powerful antibiotic using AI' and a sub-headline 'Machine learning uncovers potent new drug able to kill 35 powerful bacteria'. The date '21 February 2020' is listed next to the sub-headline. To the right of the main content, a vertical sidebar titled 'TIMES' shows a section titled 'HOW TO SPEND IT' with an image of a coronavirus and a downward-pointing arrow. Another article in the sidebar discusses 'Anti-social robots' and their impact on social distancing. The bottom right corner of the image features a small blue circular logo with a white swirl pattern.

Argua

# NEWS

Home | Video | World | UK | Business | Tech | Science | Stories | Entertainment & Arts

BBC WORKLIFE

## Our new guide for getting ahead

# Scientists discover powerful antibiotic using AI

Machine learning uncovers potent new drug able to kill 35 powerful bacteria

21 February 2020

Share

(Stokes et al., Cell'20)

Machine learning uncovers potent new drug able to kill 35 powerful bacteria

TIMES



HOW TO SPEND IT

Anti-social robots harr increase social distanc

g-resistant

© 2020

# GNNs are a very hot research topic



GNNs are currently experiencing their  
“ImageNet” moment



# Rich ecosystem of libraries



PyTorch  
geometric


[github.com/rusty1s/pytorch\\_geometric](https://github.com/rusty1s/pytorch_geometric)

DGL  
[dgl.ai](https://dgl.ai)



Spektral

*graphneural.network*



[github.com/deepmind/graph\\_nets](https://github.com/deepmind/graph_nets)



[github.com/deepmind/jraph](https://github.com/deepmind/jraph)

## Rich ecosystem of datasets




[ogb.stanford.edu](https://ogb.stanford.edu)



PyTorch  
geometric

<https://pytorch-geometric.readthedocs.io/en/latest/modules/datasets.html>



TUDataset


[graphlearning.io](http://graphlearning.io)

# Benchmarking Graph Neural Networks


[github.com/graphdeeplearning/benchmarking-gnns](https://github.com/graphdeeplearning/benchmarking-gnns)



# How to process the graph?



$$g(\mathbf{x}_b, \mathbf{X}_{\mathcal{N}_b})$$



$$\mathbf{X}_{\mathcal{N}_b} = \{\{\mathbf{x}_a, \mathbf{x}_b, \mathbf{x}_c, \mathbf{x}_d, \mathbf{x}_e\}\}$$




# What's in a GNN layer?

- We construct useful functions over graphs,  $f$ , by shared application of a local permutation-invariant function  $g(x_i, X_{N_i})$ .
  - We often refer to  $f$  as “GNN layer”,  $g$  as “diffusion”, “propagation”, “message passing”
- We will take a quick look at ways in which we can actually concretely **define**  $g$ .
  - **Very intense** area of research!
- Fortunately, *almost all* proposed layers can be classified as one of three ***spatial*** “flavours”.



# The three “flavours” of GNN layers



$$\mathbf{h}_i = \phi \left( \mathbf{x}_i, \bigoplus_{j \in \mathcal{N}_i} c_{ij} \psi(\mathbf{x}_j) \right)$$

$$\mathbf{h}_i = \phi \left( \mathbf{x}_i, \bigoplus_{j \in \mathcal{N}_i} a(\mathbf{x}_i, \mathbf{x}_j) \psi(\mathbf{x}_j) \right)$$

$$\mathbf{h}_i = \phi \left( \mathbf{x}_i, \bigoplus_{j \in \mathcal{N}_i} \psi(\mathbf{x}_i, \mathbf{x}_j) \right)$$




# Convolutional GNN

- Features of neighbours aggregated with fixed weights,  $c_{ij}$

$$\mathbf{h}_i = \phi \left( \mathbf{x}_i, \bigoplus_{j \in \mathcal{N}_i} c_{ij} \psi(\mathbf{x}_j) \right)$$

- Usually, the weights depend directly on A.
  - ChebyNet (Defferrard et al., NeurIPS'16)
  - GCN (Kipf & Welling, ICLR'17)
  - SGC (Wu et al., ICML'19)
- Useful for **homophilous** graphs and **scaling up**
  - When edges encode *label similarity*




# Attentional GNN

- Features of neighbours aggregated with **implicit** weights (via *attention*)

$$\mathbf{h}_i = \phi \left( \mathbf{x}_i, \bigoplus_{j \in \mathcal{N}_i} a(\mathbf{x}_i, \mathbf{x}_j) \psi(\mathbf{x}_j) \right)$$

- Attention weight computed as  $\alpha_{ij} = a(\mathbf{x}_i, \mathbf{x}_j)$ 
  - MoNet (Monti et al., CVPR'17)
  - GAT (Veličković et al., ICLR'18)
  - GaAN (Zhang et al., UAI'18)
- Useful as “middle ground” w.r.t. **capacity** and **scale**
  - Edges need not encode homophily
  - But still compute scalar value in each edge




# Message-passing GNN

- Compute **arbitrary vectors** (“messages”) to be sent across edges

$$\mathbf{h}_i = \phi \left( \mathbf{x}_i, \bigoplus_{j \in \mathcal{N}_i} \psi(\mathbf{x}_i, \mathbf{x}_j) \right)$$


- Messages computed as  $\mathbf{m}_{ij} = \psi(\mathbf{x}_i, \mathbf{x}_j)$ 
  - Interaction Networks (Battaglia et al., NeurIPS'16)
  - MPNN (Gilmer et al., ICML'17)
  - GraphNets (Battaglia et al., 2018)
- Most **generic** GNN layer
  - May have *scalability* or *learnability* issues
  - Ideal for *computational chemistry, reasoning and simulation*



Message-passing




# How to use GNNs?




Inputs  
( $\mathbf{X}$ ,  $\mathbf{A}$ )




# How to use GNNs?




# How to use GNNs?



# How to use GNNs?




# How to use GNNs?



# If you'd like to know more

For (substantially!) more context, I recently gave a talk on **theoretical GNN foundations**:  
<https://www.youtube.com/watch?v=uF53xsT7mjc>



## *...Back to the past*

- In 2017, as part of my Mila internship we proposed Graph Attention Networks (GATs)
  - One of the first prominent examples of attentional GNN
  - They remain a popular model to this day
- It was only loosely clear that models like these could benefit my biological projects
  - We set out to find out exactly how...





# CNNs\* for Mesh-based Parcellation of the Cerebral Cortex

Guillem Cucurull, Konrad Wagstyl, Arantxa Casanova, **Petar Veličković**,  
Estrid Jakobsen, Michal Drozdzal, Adriana Romero, Alan Evans and Yoshua Bengio




# Cortex parcellation

- Different areas of the cerebral cortex are involved in different cognitive processes
  - Visual processing
  - Language comprehension
- Mapping these areas helps us understand how the cortex is organised
- Our graph attention network paper was, in fact, built for this very purpose :)
- We focus on regions 44 and 45 of **Broca's area**:




# What is a cortical mesh?


- Common coordinate system
- Can represent multiple modalities and features
- Can be used to coregister cortical surfaces between different individuals
- We can run a GNN over the nodes in the mesh!
  - Classify nodes as “44”, “45”, or “background”.




# Quantitative results



# Quantitative results




# Quantitative results




# Quantitative results

+ Node **positional** features!





# Qualitative results




# With hindsight...

- Meshes come with a *lot* of useful geometry
  - Evident by utility of positional features
  - (Vanilla) GNNs would *discard* that information
- We now have a **wealth** of architectures that are specialising for the mesh domain!
  - Geodesic CNN (Masci *et al.*)
  - MoNet (Monti *et al.*)
  - Gauge Equivariant Mesh CNN (de Haan *et al.*)
- All of the above would make great choices for processing the brain mesh!
  - Perhaps an interesting future project? 🤔
  - Reach out to me if you're curious!



(a) Parallel transport on a flat mesh.



(b) Parallel transport along an edge of a general mesh.

## *...Back to the past* 🧠

- This project proved to me the untapped utility that GNNs can have in biological problems
  - We applied the GCN and GAT models pretty much out-of-the-box!
- Now was the time to revisit my earlier collaboration (Parapred) under this lens.




# Attentive Cross-modal Paratope Prediction

Andreea Deac, Petar Veličković and Pietro Sormanni



# Motivation for antibody design

- Antibodies are
  - Y-shaped proteins
  - a critical part of our immune system
- They neutralise pathogenic bacteria and viruses by tagging the antigen in a “lock and key” system.
- Designing our own arbitrary antibodies would be a big step towards personalised medicine.
- (You've probably heard a whole lot about antibodies and antigens in the past year...)




# Towards personalised medicine


- Generating an antibody requires first predicting the specific amino acids (**the paratope**) which participate in the neutralisation of the antigen.
- **Input:** a sequence of (one-hot encoded) antibody amino acids.  
(+ a sequence of (one-hot encoded) antigen amino acids)
- **Output:** probability for each amino acid to participate in the binding with the antigen.




# Paratope prediction




# Parapred and Fast-Parapred architecture




# Paratope prediction (+ antigen)




# Fast-Parapred




# AG-Fast-Parapred



# AG-Fast-Parapred





# Quantitative results



	ROC AUC	MCC	Epoch time
proABC	0.851	0.522	—
Parapred	$0.880 \pm 0.002$	$0.564 \pm 0.005$	$0.190 \pm 0.019\text{s}$
Fast-Parapred	$0.883 \pm 0.001$	$0.572 \pm 0.004$	$0.085 \pm 0.015\text{s}$
AG-Fast-Parapred	$0.899 \pm 0.004$	$0.598 \pm 0.012$	$0.178 \pm 0.020\text{s}$



# Qualitative results



The model learns the antibody/antigen **geometry** without being given any positional information!



## *...Back to the past*

- Now it was apparent that stitching GNNs into protein-protein interaction made sense!
- Could we explore some other cases of molecular interaction?




# Drug-Drug Adverse Effect Prediction with Graph Co-Attention

Andreea Deac, Yu-Hsiang Huang, Petar Veličković, Pietro Liò and Jian Tang



# Drug use is increasing

	2000	2011
Prescription Drug Use	51%	59%
>5 drugs	8.2%	15%



# Polypharmacy

- Polypharmacy is the concurrent use of multiple medications by a patient.
- It is necessary for chronic, complex or multiple conditions and most of the increase in cost comes from treating these.
- “Hulk & Iron Man” analogy: drugs correspond to ‘heroes’, but putting them together can destroy the surrounding city!






# Adverse side-effects

- Side effects affecting 15% of the population, treatment costs exceeding \$177 billion/year
- Some found in Phase IV of clinical trials
- But plenty are undiscovered when the drugs are put on the market



# Related work



- Most models predict if a side-effect exists or not (using drug-drug similarity: chemical substructures, individual drug side effects, interaction profile fingerprints)
- Others model the interactions between pairs of drugs, pairs of proteins and drug-protein pairs to predict “missing” links between them.
- We, instead, represents **molecules as graphs!**




\* Modeling polypharmacy side effects with graph convolutional networks, Žitnik et al, 2018



# DDI - Tasks



# Graph co-attention



**Drug *x***


$$(d_x)m_{ij}^t = f_e^t \left( e_{ij}^{(d_x)} \right) \odot f_v^t \left( (d_x)h_j^{t-1} \right)$$

**Drug *y***

$$(d_x)n_i^t = f_o^t \left( \sum_{k=1}^K \sum_{\forall j \in d_y} (k)\alpha_{ij}^t \cdot (k)\mathbf{W}_v^{t(d_y)} h_j^{t-1} \right)$$



# The (MH)CADDI Architecture



$$(d_x) h_i^t = \text{LayerNorm} \left( (d_x) h_i^{t-1} + (d_x) m_i^t + (d_x) n_i^t \right)$$

$$d_x = \sum_{\forall j \in d_x} f_r \left( (d_x) h_j^T \right)$$



## Variants considered

- *MPNN-Concat*: removing co-attention, i.e. learning drug representations independently;
- *Late-Outer*: where co-attention messages are not aggregated until the last layer;
- *CADDI*: only  $K = 1$  attention head.



# Quantitative results

**Table 1:** Comparative evaluation results after stratified 10-fold crossvalidation.

	AUROC
<b>Drug-Fingerprints</b> [21]	0.744
<b>RESCAL</b> [30]	0.693
<b>DEDICOM</b> [31]	0.705
<b>DeepWalk</b> [32]	0.761
<b>Concatenated features</b> [46]	0.793
<b>Decagon</b> [46]	0.872
<b>MHCADDI</b> (ours)	<b>0.882</b>
<b>MHCADDI-ML</b> (ours)	0.819

**Table 2:** Ablation study for various aspects of the MHCADDI model.

	AUROC
<b>MPNN-Concat</b>	0.661
<b>Late-Outer</b>	0.724
<b>CADDI</b>	0.778
<b>MHCADDI</b>	<b>0.882</b>



## *...Back to the past*

- It was ~at this point I graduated from my PhD, and joined DeepMind
- Gradually oriented back towards classical algorithms, and away from biology
  - Luckily, biology is **packed** with interesting classical algorithms :)
- The following three works (time permitting) represent a medley of biological approaches I was involved in during this time.
  - Two of these opportunities came *not far from home* :)
  - The third one was **years** in the making!



# Hierarchical Protein Function Prediction with Tail-GNNs

Stefan Spalević, **Petar Veličković**, Jovana Kovačević and Mladen Nikolić




# Protein function prediction

- Detecting mechanisms of action for proteins is a highly relevant task!
- It is also an area ripe with graphs!
  - Protein itself can be represented as a graph (if known structure; Gligorijević et al.)
  - Protein–protein interaction networks are graphs (standard **PPI** benchmark for GNNs)
  - In this particular domain, a graph comes up in one more place...



# Protein function prediction


- The label space of functions is itself a graph! (**gene ontology**)



- Requires a GNN in the **label space**
  - Our literature survey suggested no proposals like this!
  - Once again, a biological problem motivates a core architecture



# Tail-GNN



# Quantitative results


- With the right aggregator choice + spectral features, yields significant benefits!

Model	Validation $F_1$	Test $F_1$
Labelling network	$0.582 \pm 0.003$	$0.584 \pm 0.003$
Tail-GNN-mean	$0.583 \pm 0.006$	$0.586 \pm 0.004$
Tail-GNN-GAT	$0.582 \pm 0.004$	$0.587 \pm 0.005$
Tail-GNN-max	$0.581 \pm 0.002$	$0.585 \pm 0.004$
Tail-GNN-sum	<b><math>0.596 \pm 0.003</math></b>	<b><math>0.600 \pm 0.003</math></b>
Tail-GNN-sum (no spectral fts.)	$0.587 \pm 0.007$	$0.590 \pm 0.008$



# A Step Towards Neural Genome Assembly

Lovro Vrček, Petar Veličković and Mile Šikić



# Genome assembly



stack of NY Times, June 27, 2000



stack of NY Times, June 27, 2000  
on a pile of dynamite



this is just hypothetical



so, what did the June 27, 2000 NY  
Times say?



# Genome assembly

Multiple Copies of a Genome (Millions of them)




CTGATGATGGACTACGCTACTACTGCTAGCTGTATTACGATCAGCTACCACATCGTAGCTACGATGCACTAGGATCAGCTACCATCGTACG  
CTGATGATGGACTACGCTACTACTGCTAGCTGTATTACGATCAGCTACCACATCGTAGCTACGATGCACTAGGATCAGCTACCATCGTACG  
CTGATGATGGACTACGCTACTACTGCTAGCTGTATTACGATCAGCTACCACATCGTAGCTACGATGCACTAGGATCAGCTACCATCGTACG  
CTGATGATGGACTACGCTACTACTGCTAGCTGTATTACGATCAGCTACCACATCGTAGCTACGATGCACTAGGATCAGCTACCATCGTACG

Breaking the Genomes at Random Positions



CTGATGATGGACTACGCTACTACTGCTGTATTACGATCAGCTACCACATCGTAGCTACGATGCACTAGGATCAGCTACCATCGTACG  
CTGATGATGGACTACGCTACTACTGCTGTATTACGATCAGCTACCACATCGTAGCTACGATGCACTAGGATCAGCTACCATCGTACG  
CTGATGATGGACTACGCTACTACTGCTGTATTACGATCAGCTACCACATCGTAGCTACGATGCACTAGGATCAGCTACCATCGTACG  
CTGATGATGGACTACGCTACTACTGCTGTATTACGATCAGCTACCACATCGTAGCTACGATGCACTAGGATCAGCTACCATCGTACG

“Burning” Some Reads



CTGATGATGGACTACGCTACTACTGCTGTATTACGATCAGCTACCACATCGTAGCTACGATGCACTAGGATCAGCTACCATCGTACG  
CTGATGATGGACTACGCTACTACTGCTGTATTACGATCAGCTACCACATCGTAGCTACGATGCACTAGGATCAGCTACCATCGTACG  
CTGATGATGGACTACGCTACTACTGCTGTATTACGATCAGCTACCACATCGTAGCTACGATGCACTAGGATCAGCTACCATCGTACG  
CTGATGATGGACTACGCTACTACTGCTGTATTACGATCAGCTACCACATCGTAGCTACGATGCACTAGGATCAGCTACCATCGTACG




# Genome assembly

ATCAGCTACCA  
TACTGCTAG  
CTGATGATGGACT  
ATCAGCTACC  
TGGACTACGCTAC  
AGCTATCGG  
AGCTACGATGCA  
ATGCATTAGCA  
CATCGTAGC  
CTGATGA  
TACTGTAGCT  
ATCAGCTACC  
TTAGCAAGCT  
TCGTAGCTAG  
CTGATGATGG  
TCAGCTACCA  
ATGCATTAGCAA  
ATGCATTAGCA  
CACATCGTAGC  
TACCACATCGT  
CTGATGATGG  
ATCGTAGCTACG  
TACTGCTAGCT  
ACATCGTAGCT  
ATCGGATCA  
GGATCAGCTAC  
ACTACTGCTA  
TACGCTAC  
GCAAGCTATC  
ACTACGCTAC  
AGCTAC  
GTATTACGATC  
GCTAGCTGTAT



# Genome assembly using Hamiltonian paths



# But... the reads are **faulty**!

- Learn **algorithms** to prune errors




Figure 1: Example of structures in the assembly graph, before all the simplification steps. Letter **A** marks transitive edges, a short tip is marked with **B**, and a bubble which cannot be fully resolved is marked with **C**. Red crosses show which edges can be removed from the assembly graph.



# Towards neural genome assembly

Pre-train on **synthetic** graphs...

...generalises to **real** organisms!

(Still preliminary, but **encouraging!**)

Table 1: Scaling of algorithm execution for isolated learning of algorithms.

Algorithm	Scaling				
	1x	2x	4x	8x	20x
Transitive removal	98.10%	99.00%	99.52%	99.76%	99.91%
Tips trimming	98.05%	98.96%	99.49%	99.70%	99.87%
Bubble popping	98.16%	99.03%	99.53%	99.77%	99.90%

Table 2: Scaling of algorithm execution for parallel learning of algorithms.

Algorithm	Scaling				
	1x	2x	4x	8x	20x
Transitive removal	98.21%	99.07%	99.50%	99.89%	99.92%
Tips trimming	98.45%	99.11%	99.46%	99.76%	99.89%
Bubble popping	98.17%	99.02%	99.51%	99.78%	99.90%

Table 3: Parallel algorithm execution on the assembly graph of lambda phage.

	Transitive removal	Tips trimming	Bubble popping
Lambda phage	98.04%	93.33%	97.47%
E. coli	99.67%	98.84%	99.26%



# Further insight: Algorithmic reasoning

If you would like to know more details about teaching GNNs to be more “algorithmic”:



<https://www.youtube.com/watch?v=IPO6CPoluok>



[https://drive.google.com/file/d/1\\_EQ9Yu7VEkvrHaVH1\\_WbT5ABvxrSNY-s/view?usp=sharing](https://drive.google.com/file/d/1_EQ9Yu7VEkvrHaVH1_WbT5ABvxrSNY-s/view?usp=sharing)



# Broader context: Combinatorial Optimisation

## Combinatorial optimization and reasoning with graph neural networks

Quentin Cappart<sup>1</sup>, Didier Chételat<sup>2</sup>, Elias Khalil<sup>3</sup>, Andrea Lodi<sup>2</sup>,  
Christopher Morris<sup>2</sup>, and Petar Veličković<sup>\*4</sup>

<sup>1</sup>Department of Computer Engineering and Software Engineering, Polytechnique Montréal

<sup>2</sup>CERC in Data Science for Real-Time Decision-Making, Polytechnique Montréal

<sup>3</sup>Department of Mechanical & Industrial Engineering, University of Toronto

<sup>4</sup>DeepMind

Our 43-page survey on GNNs for CO!

<https://arxiv.org/abs/2102.09544>

Section 3.3. details algorithmic reasoning,  
with comprehensive references.

Combinatorial optimization is a well-established area in operations research and computer science. Until recently, its methods have focused on solving problem instances in isolation, ignoring the fact that they often stem from related data distributions in practice. However, recent years have seen a surge of interest in using machine learning, especially graph neural networks (GNNs), as a key building block for combinatorial tasks, either as solvers or as helper functions. GNNs are an inductive bias that effectively encodes combinatorial and relational input due to their permutation-invariance and sparsity awareness. This paper presents a conceptual review of recent key advancements in this emerging field, aiming at both the optimization and machine learning researcher.




# Predicting Patient Outcomes with Graph Representation Learning

Emma Rocheteau\*, Catherine Tong\*, Petar Veličković, Nicholas Lane and Pietro Liò




# Electronic Health Records (EHRs) in the ICU

- EHRs can provide plentiful information about a patient's progression
  - But not all data contained in there are easy to leverage by deep learning systems!
- Today, we focus on **diagnoses**.



# Diagnosis information is hard to use


- Large number of possibilities makes distinguishing *patterns* of comorbidity **difficult**.



- There is a lack of data for rarer combinations.
  - A long tail of **rare** diagnoses, difficult for deep learning models to leverage!



# Distribution of diagnoses in eICU



# The “pattern recognition” method


- Commonly, the “long tail” of diagnoses is *discarded* and the rest embedded.
  - But this long tail often holds the most **useful** cues, which diagnosticians regularly use!
- How do **clinicians** often make decisions about diagnoses or prognoses?
- The *pattern recognition* diagnostic method, as described by Wikipedia:

*“In a pattern recognition method the provider uses **experience** to recognize a pattern of clinical characteristics... This may be the primary method used in cases where diseases are “obvious”, or the provider’s **experience** may enable him or her to **recognize** the condition quickly.”*

- We interpret experience as exploitation of related cases the clinician treated in the past.
  - Hence, the cases form a **graph**!




# These links definitely exist :)



# The graph of patients

- Key assumption: patients with **related diagnoses** will likely have **related prognoses**!



- If we use this signal wisely, it can be a great way to regularise our model **and** make advantage of sparse diagnosis data.



# How to build the graph?

- The “relatedness” score between two patients  $i$  and  $j$  is given by:

$$\mathcal{M}_{ij} = \alpha \sum_{\mu=1}^m \mathcal{D}_{i\mu} \mathcal{D}_{j\mu} (d_\mu^{-1} + \gamma) - \sum_{\mu=1}^m \mathcal{D}_{i\mu} + \mathcal{D}_{j\mu}$$


where:

- $\mathbf{D}$  is a *diagnosis matrix* (s.t.  $D_{i\mu}$  means “does patient  $i$  have diagnosis  $\mu$ ?“)
  - $d_\mu$  is the *frequency* of diagnosis  $\mu$
  - $m$  is the *number* of diagnoses
  - $\alpha$  and  $\gamma$  are *hyperparameters*
- 
- Can **threshold** based on the relatedness scores




# Hybrid LSTM-GNN model


① Construct Graph      ② LSTM Embeddings      ③ GNN Embeddings




④ Static Embeddings




⑤ Concatenate



# Our results



# Qualitative: LSTM-GAT Attention weights



# AAAI'21 Workshop Recognition

## Awards

Best short paper (\$250 winner, \$125 runner-up)

### Runner-up

Emma Rocheteau, Catherine Tong, Petar Veličković, Nicholas Lane and Pietro Liò. *Predicting Patient Outcomes with Graph Representation Learning*

### Winner

Beatrice Portelli, Daniele Passabì, Edoardo Lenzi, Giuseppe Serra, Enrico Santus and Emmanuele Chersoni. *Improving Disease Drug Event Extraction with SpanBERT on Different Text Types*



# In conclusion...

- Studying biological problems with graph representation learning is likely here to stay
  - Abundance of data “sitting and waiting to be processed”
  - In many problems of interest, state-of-the-art is still a **shallow** method
  - Often, biological problems can give rise to **core** methodological progress.
- With the right mindset, no proper biological training is needed!
  - Just the ability to carefully listen, and work **together** with biologists.
- For biologists: I hope I've convinced you that GNNs could be a useful tool!
- But ultimately, I would love to stimulate, and see even more of, **interdisciplinary** research.



# Thank you!

Questions?

[petarv@google.com](mailto:petarv@google.com) | <https://petar-v.com>

With many thanks to:

**Lead authors:** Guillem Cucurull, Andreea Deac, Stefan Spalević, Lovro Vrček, Emma Rocheteau, Catherine Tong  
**Bio/chem:** Konrad Wagstyl, Estrid Jakobsen, Alan Evans, Pietro Sormanni, Pietro Lio, Mile Šikić, Jovana Kovačević

**ML collaborators:** Arantxa Casanova, Yu-Hsiang Huang

**ML advisors:** Adriana Romero, Michal Drozdzal, Yoshua Bengio, Jian Tang, Mladen Nikolić, Nic Lane

